

P8106_HW4

Naomi Simon-Kumar

ns3782

04/11/2025

Contents

Loading libraries	1
Part 1: Tree Based Models using College Data	2
Partition into training and testing set	2
a) Build a regression tree on training data	3
b) Perform random forest on training data	8
c) Perform boosting on training data	12

Loading libraries

```
# Load libraries
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.1      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.1
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(tidymodels)
```

```
## -- Attaching packages ----- tidymodels 1.2.0 --
## v broom       1.0.7      v rsample    1.2.1
## v dials        1.3.0      v tune       1.2.1
## v infer        1.0.7      v workflows  1.1.4
## v modeldata    1.4.0      v workflowsets 1.1.0
## v parsnip      1.2.1      v yardstick  1.3.2
## v recipes      1.1.0
```

```
## -- Conflicts ----- tidymodels_conflicts() --
## x scales::discard() masks purrr::discard()
## x dplyr::filter() masks stats::filter()
## x recipes::fixed() masks stringr::fixed()
## x dplyr::lag() masks stats::lag()
## x yardstick::spec() masks readr::spec()
## x recipes::step() masks stats::step()
## * Dig deeper into tidy modeling with R at https://www.tmwr.org
```

```
library(caret)
```

```
## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following objects are masked from 'package:yardstick':
##
##   precision, recall, sensitivity, specificity
##
## The following object is masked from 'package:purrr':
##
##   lift
```

```
library(ggplot2)
library(rpart)
```

```
##
## Attaching package: 'rpart'
##
## The following object is masked from 'package:dials':
##
##   prune
```

```
library(rpart.plot)
library(ranger)
```

```
## Warning: package 'ranger' was built under R version 4.4.3
```

Part 1: Tree Based Models using College Data

Partition into training and testing set

```
# Read in dataset
college <- read.csv("College.csv")

# Remove NAs
college <- na.omit(college)

# Set seed for reproducibility
```

```

set.seed(299)

# Split data into training and testing data
data_split_college <- initial_split(college, prop = 0.8)

# Extract the training and test data, removing college ID column
training_data_college <- training(data_split_college) %>% select(-College)
testing_data_college <- testing(data_split_college) %>% select(-College)

```

a) Build a regression tree on training data

In order to implement the CART approach implementing recursive partitioning and pruning, I first fit a regression tree using `cp=0` (complexity parameter). This parameter controls the complexity pruning in the CART algorithm, i.e., how splits are undertaken. Setting `cp=0` was a safe choice to ensure that the tree was sufficiently large, allowing all potential splits are considered. I also produced a plot of this tree.

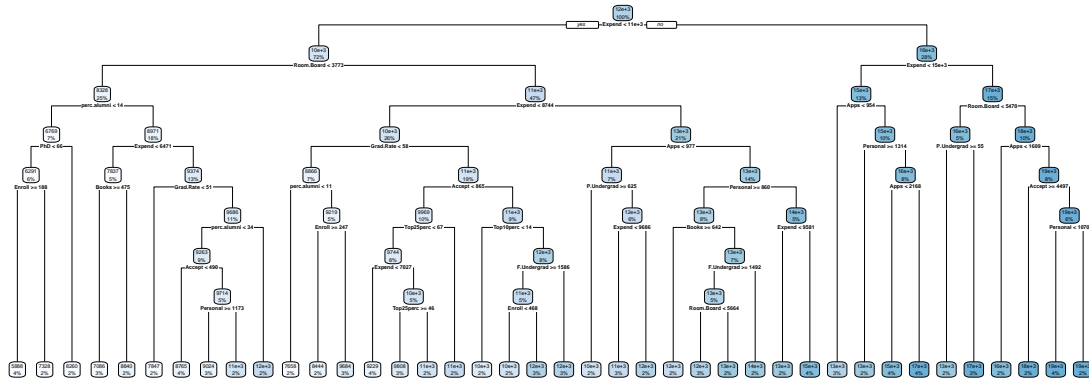
```

# Set seed for reproducibility
set.seed(299)

# Fit initial tree:
initial.tree.fit <- rpart(Outstate ~ ., data = training_data_college,
                        control = rpart.control(cp = 0))

# Tree plot
rpart.plot(initial.tree.fit)

```




```
# Print and plot the cp table
printcp(initial.tree.fit)
```

```
##
## Regression tree:
## rpart(formula = Outstate ~ ., data = training_data_college, control = rpart.control(cp = 0))
##
## Variables actually used in tree construction:
## [1] Accept      Apps        Books        Enroll       Expend      F.Undergrad
## [7] Grad.Rate   P.Undergrad perc.alumni Personal      PhD         Room.Board
## [13] Top10perc   Top25perc
##
## Root node error: 6222135701/452 = 13765787
##
## n= 452
##
##      CP nsplit rel error  xerror   xstd
## 1  0.49965968      0  1.00000 1.00404 0.058031
## 2  0.10266905      1  0.50034 0.54145 0.034753
## 3  0.05071124      2  0.39767 0.44785 0.030837
## 4  0.04021488      3  0.34696 0.43401 0.032119
## 5  0.01820952      4  0.30675 0.39616 0.031612
## 6  0.01319074      5  0.28854 0.37639 0.032586
## 7  0.01255106      7  0.26215 0.37556 0.032802
## 8  0.01194445      8  0.24960 0.37611 0.032884
```

## 9	0.01165684	9	0.23766	0.37553	0.033331
## 10	0.00942611	10	0.22600	0.37590	0.033641
## 11	0.00770392	11	0.21658	0.36867	0.032816
## 12	0.00672841	12	0.20887	0.37730	0.033123
## 13	0.00619338	13	0.20214	0.38085	0.033378
## 14	0.00587852	14	0.19595	0.38098	0.033357
## 15	0.00538113	15	0.19007	0.37111	0.032927
## 16	0.00531018	17	0.17931	0.37395	0.033274
## 17	0.00377286	18	0.17400	0.37834	0.033610
## 18	0.00353985	19	0.17023	0.37250	0.032941
## 19	0.00254368	20	0.16669	0.37870	0.033291
## 20	0.00246805	21	0.16414	0.38447	0.033757
## 21	0.00212300	22	0.16167	0.38576	0.033736
## 22	0.00210237	23	0.15955	0.38858	0.033838
## 23	0.00207706	24	0.15745	0.39037	0.033877
## 24	0.00199280	25	0.15537	0.38740	0.033677
## 25	0.00188271	26	0.15338	0.38777	0.034094
## 26	0.00187160	27	0.15150	0.38754	0.034083
## 27	0.00179477	28	0.14963	0.38986	0.034280
## 28	0.00167899	30	0.14604	0.39016	0.034537
## 29	0.00139108	31	0.14436	0.38893	0.034479
## 30	0.00120513	32	0.14297	0.38988	0.033974
## 31	0.00119267	33	0.14176	0.39068	0.033937
## 32	0.00110541	34	0.14057	0.39051	0.033940
## 33	0.00099926	36	0.13836	0.39288	0.033968
## 34	0.00073074	37	0.13736	0.39417	0.033983
## 35	0.00031851	38	0.13663	0.39388	0.033991
## 36	0.00000000	39	0.13631	0.39305	0.033715

```
plotcp(initial.tree.fit)
```




To explore the impact of adjusting the complexity parameter, I also fit a second model setting $cp=0.01$. This model had fewer splits by comparison. When plotted, this tree was noticeably smaller than the previous tree.

```
# Set seed for reproducibility
set.seed(299)

# Fit another tree:
tree.fit.2 <- rpart(Outstate ~ ., data = training_data_college,
                   control = rpart.control(cp = 0.01))


# Tree plot
rpart.plot(tree.fit.2)
```



```
# Print and plot the cp table
printcp(tree.fit.2)
```

```
##
## Regression tree:
## rpart(formula = Outstate ~ ., data = training_data_college, control = rpart.control(cp = 0.01))
##
## Variables actually used in tree construction:
## [1] Apps      Expend    Grad.Rate P.Undergrad perc.alumni Room.Board
##
## Root node error: 6222135701/452 = 13765787
##
## n= 452
##
##      CP nsplit rel error  xerror   xstd
## 1  0.499660     0  1.00000  1.00404  0.058031
## 2  0.102669     1  0.50034  0.54145  0.034753
## 3  0.050711     2  0.39767  0.44785  0.030837
## 4  0.040215     3  0.34696  0.43401  0.032119
## 5  0.018210     4  0.30675  0.39616  0.031612
## 6  0.013191     5  0.28854  0.37639  0.032586
## 7  0.012551     7  0.26215  0.37556  0.032802
## 8  0.011944     8  0.24960  0.37611  0.032884
## 9  0.011657     9  0.23766  0.37553  0.033331
## 10 0.010000    10  0.22600  0.37024  0.033137
```

```
plotcp(tree.fit.2)
```



Importantly, setting $cp=0$ is the preferred choice, as it allowed for a large enough tree to be grown for the cost complexity table. In the $cp = 0.01$ model, the smallest error (scaled cross-validation error) was 0.37369, whereas the $cp = 0$ model achieved a slightly lower minimum error of 0.36867. This shows us that a fully grown tree in this case is better suited for selecting an optimal complexity parameter based on cross-validation.

The optimal tree selected from the $cp = 0$ model has a **complexity parameter of 0.00770392**, with **11 splits** (i.e., 12 terminal nodes). This was the model that minimised scaled cross-validation error (xerror = 0.36867). Therefore, this was chosen as the final pruned tree.

b) Perform random forest on training data

I decided to explore two tuning grids to find the optimal random forest model using cross-validation RMSE. The `mtry` parameter controls the number of predictors randomly selected at each split in the forest. I decided to tune `mtry` over the full range of possible values, from 1 to 16 (i.e., the full number of predictors in the college dataset).

For the first grid (`min.node.size = 1:7`), the best model had `mtry = 7` and `min.node.size = 5`.

```
# Set seed for reproducibility
set.seed(299)

# Set cross-validation
ctrl <- trainControl(method = "cv")
```



```

# Define grid for tuning mtry and min.node.size
rf.grid <- expand.grid(
  mtry = 1:16, # max no. of predictors
  splitrule = "variance",
  min.node.size = c(1:7)
)

```

```

# Fit random forest using ranger via caret
rf.fit <- train(
  Outstate ~ .,
  data = training_data_college,
  method = "ranger",
  tuneGrid = rf.grid,
  trControl = ctrl
)

```

```
## Growing trees.. Progress: 54%. Estimated remaining time: 1 minute, 44 seconds.
```

```

# Obtain optimal tuning parameters from cross-validation
rf.fit$bestTune # mtry = 7, min.node.size = 5

```

```

##      mtry splitrule min.node.size
## 47      7  variance                5

```

```
ggplot(rf.fit, highlight = TRUE)
```

```


## Warning: The shape palette can deal with a maximum of 6 discrete values because more
## than 6 becomes difficult to discriminate
## i you have requested 7 values. Consider specifying shapes manually if you need
## that many have them.

```

```

## Warning: Removed 16 rows containing missing values or values outside the scale range
## (`geom_point()`).

```



I then decided to try an expanded grid range for min.node.size, extending it to 1:10 to check for any better performing values beyond the original range.

```
# Set seed for reproducibility
set.seed(299)

# Define another grid for tuning mtry and min.node.size
rf.grid2 <- expand.grid(
  mtry = 1:16, # max no. of predictors
  splitrule = "variance",
  min.node.size = c(1:10)
)

# Fit random forest using ranger via caret
rf.fit2 <- train(
  Outstate ~ .,
  data = training_data_college,
  method = "ranger",
  tuneGrid = rf.grid2,
  trControl = ctrl
)
```

```
## Growing trees.. Progress: 48%. Estimated remaining time: 12 minutes, 19 seconds.
```

```
# Optimal parameters
rf.fit2$bestTune # mtry = 9, min.node.size = 3 (row 83)
```

```
##      mtry splitrule min.node.size
## 83      9  variance                3
```


```
rf.fit2$results[83, ] # pulling the lowest RMSE: 1763.865
```

```
##      mtry splitrule min.node.size      RMSE Rsquared      MAE      RMSESD RsquaredSD
## 83      9  variance                3 1763.865 0.774242 1332.22 307.1101 0.08056556
##      MAESD
## 83 176.1794
```

```
# Plot performance for tuning grid values
ggplot(rf.fit2, highlight = TRUE)
```

```
## Warning: The shape palette can deal with a maximum of 6 discrete values because more
## than 6 becomes difficult to discriminate
## i you have requested 10 values. Consider specifying shapes manually if you need
## that many have them.
```

```
## Warning: Removed 64 rows containing missing values or values outside the scale range
## (`geom_point()`).
```



From this grid, **the optimal tuning parameters were `mtry = 9` and `min.node.size = 3`**, selected based on the lowest cross validation RMSE (1763.865). The corresponding plot shows the lowest RMSE.

Next, reporting the variable importance and the test error for this selected model (`rf.fit2`):

c) Perform boosting on training data